

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

Senior Data Engineer

Senior data engineer with 10+ years building production data pipelines at scale — ingesting and normalizing data from 220+ distinct sources into centralized analytics platforms, designing a shared data dictionary (Common Information Model) to standardize schemas across all of it, and running the full pipeline through a GitLab CI/CD process. Deep hands-on experience with Python/SQL, cloud-scale distributed compute (GCP Dataproc, BigQuery, Dataflow), and building data-quality monitoring so downstream products and features can trust what's flowing through the pipeline.

CORE SKILLS

Data Engineering & Pipelines: ETL/ELT pipeline design and ingestion for 220+ distinct data sources, Apache Beam / GCP Dataflow for large-scale historical/cold-storage data retrieval, data-quality monitoring and alerting, Common Information Model (data dictionary) design and standardization

Languages & Query: Python, SQL, PySpark / SparkSQL

Cloud & Distributed Compute: GCP (Dataproc, BigQuery, Dataflow), AWS, Azure

CI/CD & Tooling: GitLab CI/CD pipelines, Git, Docker, Zeppelin notebooks

Data Science / Analytics: Pandas, scikit-learn, NumPy, SciPy, clustering / unsupervised ML, time-series analysis

PROFESSIONAL EXPERIENCE

Senior Data/Security Engineer — Shorepoint

Oct 2023 – Present

- Built an entirely new data ingestion platform (DOE/NNSA project, completed) pulling multiple large third-party data feeds into a central Elasticsearch environment — custom data transforms, pipeline monitoring, and data-quality alerting on top of the ingested data.
- Supported data ingestion and data-quality engineering (CISA CDM project, completed) across a combined Elasticsearch/Splunk environment for a federal agency data platform.
- Currently build and maintain analytics content against a large-scale Splunk data environment for Treasury's SOC (current project).

Senior Data Scientist / Engineer — Forescout

Aug 2022 – Oct 2023

- Senior member of a data engineering/data science team building pipelines and analytics content against massive-scale customer data on cloud-based big data tooling.

Data Engineer / Data Scientist — Trend Micro / Cysiv

Sep 2018 – Aug 2022

- Owned data engineering/pipelining for 220+ unique data sources feeding a next-gen cloud-based analytics platform, from early-startup buildout through scale.
- Built 50+ Logstash filters for parsing/normalization, deployed Elasticsearch Beats for collection, and authored a Common Information Model — a data dictionary standardizing field names/types across every parsed source. Co-built "Loggify," a homegrown log parsing/filtering tool that replaced Logstash in production.
- Built a homegrown Apache Beam program run via GCP Dataflow to fetch large volumes of historical cold-storage data for customers; owned connector/collector health monitoring and pipeline troubleshooting.
- Ran exploratory data analysis at scale on GCP Dataproc compute clusters using Zeppelin notebooks and PySpark/SparkSQL; built time-series and clustering-based analytics on top of the pipeline output.

Data Scientist — Experian

Jan 2015 – Jan 2018

- Analyzed large-scale datasets to build custom models identifying anomalous patterns; built DNS-based detection/mitigation logic on top of a production data pipeline.

EDUCATION & CERTIFICATIONS

M.S. Physics — University of North Texas (2013) | B.S. Physics — Ball State University (2011)

Splunk User Certification · Splunk for Analytics and Data Science · Johns Hopkins (Coursera): R Programming, The Data Scientist's Toolbox

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

July 27, 2026

Hiring Team
flexEngage

Being the first full-time data engineer at flexEngage means setting the technology standards a whole team will build on later — that's the exact position I was in early at Cysiv, and it's the kind of problem I do my best work on.

At Trend Micro/Cysiv I owned data engineering for 220+ distinct, messy third-party data sources feeding a production analytics platform — writing the parsing/normalization layer (50+ Logstash filters), authoring a Common Information Model to give every downstream team one standardized data dictionary instead of 220 different schemas, and building a homegrown Apache Beam pipeline on GCP Dataflow to retrieve historical data at scale. That's the same shape of problem as building flexEngage's data warehouse and processing pipelines from a standing start: get the ingestion and standardization right so every feature built on top of it can trust the data.

I've also run production pipelines through a full CI/CD process (GitLab), written the exploratory analysis layer on top of pipeline output using PySpark/SparkSQL on GCP Dataproc, and worked daily in Python and SQL against large-scale cloud data platforms (BigQuery, Dataflow, plus general AWS cloud experience).

Two honest gaps against your stack: I haven't used Airflow or Snowflake specifically, and my infrastructure-as-code experience is CI/CD-centric (GitLab) rather than deep Terraform/CloudFormation/Ansible. The underlying skills transfer directly — pipeline orchestration, data-warehouse modeling, and provisioning-as-code concepts — and I'd expect to be productive on your specific toolchain quickly.

I'd welcome the chance to talk through how that background fits what you're building next.

Kyle Main