

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

July 27, 2026

Uppercase Research Hiring Team
Google

Turning static detection workflows into self-correcting, LLM-driven pipelines inside Google SecOps is exactly the problem I've spent the last several years solving from the ground up.

At Trend Micro/Cysiv, I built and led detection-as-code orchestration across nine SIEM and EDR platforms — including Google SecOps (Chronicle) itself — through each platform's native API, run end-to-end via a GitLab CI/CD pipeline. On top of that, I developed reusable GenAI-powered tooling for detection engineers: prompt engineering to identify false positives and generate new detection logic, and GenAI "skills" that automatically convert detection rules between SIEM syntaxes — functionally the same problem Uppercase is solving with autonomous rule-writing and tuning agents for YARA-L.

More recently at Shorepoint, I built an entirely new security data platform for DOE/NNSA — from raw ingestion of CrowdStrike, Suricata, and Zeek data through a full UEBA detection layer — and I currently own the core detection and alerting content for Treasury's SOC. The throughline across both is the same: design the data plumbing and the detection logic on top of it, then automate both to scale faster than the threat landscape.

One honest gap: my depth is in building and automating detection content rather than formal security-design-review or system threat-modeling engagements. But the operational discipline of running detection-as-code at scale across nine platforms translates directly to the rigor Uppercase needs in validating agentic detection output before it ships.

I'd welcome the chance to talk through how that background applies to Uppercase's roadmap.

Kyle Main