Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

August 1, 2026

Threat Research and Detection Engineering Team
Elastic

Elasticsearch has been the core detection and analytics platform across nearly my entire career — not a peripheral tool, but the system I've built from the ground up. At DOE/NNSA I built an entirely new Elasticsearch-based security data platform from scratch, including its full UEBA detection layer; at Trend Micro/Cysiv I wrote and tuned Elasticsearch Query DSL and native detection rules as core content for an ES/Kibana-based next-gen SIEM; and at DOE I supported data ingestion and data-quality work inside a combined Elasticsearch/Splunk environment. Few candidates bring that kind of direct, cross-employer, full-lifecycle Elasticsearch depth to a Principal Security ML Research Engineer role.

Beyond Elasticsearch, my work has centered on ML-based threat detection: clustering devices on the network by behavioral features, time-series anomaly detection across authentication and process-chain behaviors, and building signature, statistical, behavioral, and ML detection rules with formally tracked false-positive rates and staged rollout before full production deployment. At Trend Micro/Cysiv, this work scaled to 2,300+ detection rules across most of the MITRE ATT&CK matrix, engineered to catch real threats while minimizing false positives for analysts downstream.

I've also put LLM APIs into production for security use cases directly — prompt engineering for detection triage and rule generation, GenAI-driven orchestration of SIEM APIs across nine platforms, and GenAI-powered tooling that converts detection rules between SIEM syntaxes for other engineers. That combination of deep Elasticsearch expertise, ML detection engineering, and hands-on production LLM integration maps directly to what this role is asking for.

I'd welcome the chance to talk through how this background could contribute to Elastic Security's threat research roadmap.

Kyle Main