

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

August 1, 2026

Data Engineering Hiring Team

TRM Labs

A lakehouse is only as trustworthy as the metadata and schema discipline underneath it. That's the exact problem I solved as an early hire at a fast-growing security data platform: I built the data engineering layer for 220+ distinct, messy sources from the ground up, including a Common Information Model that standardized field names and types across every one of them — giving every downstream consumer one consistent schema to build on instead of 220 different ones.

That ground-up ownership carries directly into the stack this role runs on. I've built and operated production pipelines on GCP end to end — Dataproc compute clusters for large-scale exploratory analysis, PySpark/SparkSQL for distributed processing, GCP Dataflow (via a homegrown Apache Beam program) for high-volume historical and cold-storage retrieval, and GCP serverless functions for event-driven enrichment. I've also worked with Kafka, and I run every pipeline change through a full CI/CD process with staged rollout and data-quality monitoring before it hits production — the same operational discipline a petabyte-scale lakehouse demands.

I'd welcome the chance to bring that pipeline-ownership and metadata-management background to TRM Labs as you scale the lakehouse architecture powering your analytics platform.

Kyle Main