Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

Senior Data Engineer — Cloud Data Pipelines & Lakehouse Infrastructure

Data engineer with 12 years building and operating large-scale data pipelines and analytics platforms, with deep hands-on GCP experience (Dataproc, BigQuery, Dataflow) and distributed processing via Spark/PySpark/SparkSQL. Built and scaled ingestion for 220+ heterogeneous data sources and designed a Common Information Model — a company-wide data dictionary standardizing schema and metadata across all of it — the same discipline modern lakehouse metadata/catalog management requires. Runs production pipelines through a full CI/CD process end to end, from ingestion through transformation to data-quality monitoring and alerting.

CORE SKILLS

Cloud Data Platform (GCP): Dataproc, BigQuery, Dataflow, GCP serverless/event-driven enrichment pipelines, Apache Beam for large-scale historical/cold-storage data retrieval

Distributed Processing & Streaming: Spark / PySpark / SparkSQL for large-scale distributed analytical processing; Kafka

Data Modeling & Metadata Management: Designed and built a Common Information Model (CIM) — a data dictionary standardizing field names/types across 220+ ingested sources — direct, hands-on schema-governance and metadata-management discipline

Data Pipeline Engineering: ETL/ELT pipeline design and ingestion at scale (220+ distinct sources), 50+ Logstash filters for parsing/normalization, data-quality monitoring and alerting, staged/safe rollout for production deployments

Orchestration, CI/CD & Languages: Full CI/CD pipeline orchestration (GitLab) for production data pipeline deployment; Python, SQL

PROFESSIONAL EXPERIENCE

Senior Data/Security Engineer — Shorepoint *Oct 2023 – Present*

- Built an entirely new data ingestion platform (DOE/NNSA project, completed) pulling multiple large third-party data feeds into a central Elasticsearch environment — custom data transforms, pipeline monitoring, and data-quality alerting built on top of it.
- Supported data ingestion and data-quality engineering (CISA CDM project, completed) across a combined Elasticsearch/Splunk environment for a large agency data platform.
- Currently builds and maintains analytics content against a large-scale Splunk data environment for Treasury's SOC (current project).

Senior Data Scientist / Engineer — Forescout *Aug 2022 – Oct 2023*

- Senior member of a data engineering/data science team building pipelines and analytics content against massive-scale customer data using cloud-based big-data tooling.

Data Engineer / Data Scientist — Trend Micro / Cysiv *Sep 2018 – Aug 2022*

- Owned data engineering/pipelining for 220+ unique data sources feeding a next-gen cloud-based analytics platform, from early-startup buildout through scale.
- Built 50+ Logstash filters for parsing/normalization and authored a Common Information Model — a data dictionary standardizing field names/types across every parsed source — the single schema of record the rest of the platform was built on.
- Built a homegrown Apache Beam program run via GCP Dataflow to fetch large volumes of historical cold-storage data on request; owned pipeline/connector health monitoring and troubleshooting.
- Ran exploratory data analysis at scale on GCP Dataproc compute clusters using Zeppelin notebooks and PySpark/SparkSQL to load and process data from cloud storage buckets.

Security Data Scientist — Experian *Jan 2015 – Jan 2018*

- Analyzed large-scale datasets to build custom models identifying anomalous patterns; built DNS-based detection/mitigation logic on top of a production data pipeline.

EDUCATION & CERTIFICATIONS

M.S. Physics — University of North Texas (2013) | B.S. Physics — Ball State University (2011)

Splunk User Certification · Splunk for Analytics and Data Science · Johns Hopkins (Coursera): R Programming, The Data Scientist's Toolbox

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

August 1, 2026

Data Engineering Hiring Team
TRM Labs

A lakehouse is only as trustworthy as the metadata and schema discipline underneath it. That's the exact problem I solved as an early hire at a fast-growing security data platform: I built the data engineering layer for 220+ distinct, messy sources from the ground up, including a Common Information Model that standardized field names and types across every one of them — giving every downstream consumer one consistent schema to build on instead of 220 different ones.

That ground-up ownership carries directly into the stack this role runs on. I've built and operated production pipelines on GCP end to end — Dataproc compute clusters for large-scale exploratory analysis, PySpark/SparkSQL for distributed processing, GCP Dataflow (via a homegrown Apache Beam program) for high-volume historical and cold-storage retrieval, and GCP serverless functions for event-driven enrichment. I've also worked with Kafka, and I run every pipeline change through a full CI/CD process with staged rollout and data-quality monitoring before it hits production — the same operational discipline a petabyte-scale lakehouse demands.

I'd welcome the chance to bring that pipeline-ownership and metadata-management background to TRM Labs as you scale the lakehouse architecture powering your analytics platform.

Kyle Main