

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

August 3, 2026

Hiring Team

Anthropic — Safeguards

Robust data infrastructure is the thing that makes safety work possible at scale, and building that infrastructure is what I've spent 12 years doing — most recently underneath detection and abuse-monitoring systems that had to be right, fast, and auditable. I'd like to bring that same foundation to the Safeguards team's data pipelines, warehousing, and analytical tooling.

At Trend Micro/Cysiv, I owned the data engineering for a next-gen cloud SIEM from its earliest days as a startup — ETL/ELT pipelines for 220+ ingested log sources, 50+ Logstash parsing/normalization filters, and a Common Information Model that standardized field names and types across every source so downstream detection and analytics could trust the data. I built a homegrown Apache Beam program run via GCP Dataflow to retrieve large volumes of historical cold-storage data on demand, and ran exploratory data analysis at scale on GCP Dataproc clusters using PySpark and SparkSQL. On top of that foundation, I created and managed 2,300+ detection rules spanning most of the MITRE ATT&CK matrix — signature, statistical, behavioral, and ML-based logic — and integrated threat intelligence directly into that detection content and alert enrichment workflow, the same kind of abuse-pattern-detection work this role's data needs to support.

At Shorepoint, I built DOE/NNSA's security data integration platform from scratch — ingesting CrowdStrike, Suricata, and Zeek telemetry into a new Elasticsearch environment, then layering a UEBA detection model, data-quality monitoring and alerting, and custom Kibana dashboards on top so stakeholders could see model behavior and anomalies without needing to query the raw data themselves. That combination — pipeline ownership, data modeling, dashboarding, and detection content built on a foundation of Python, SQL, and cloud data platforms (BigQuery, Dataproc, Dataflow) — is exactly the range this Safeguards role asks for.

I'd welcome the chance to talk through how that background applies to giving Anthropic's Safeguards team the data foundation it needs to detect misuse and keep model behavior safe at scale.

Kyle Main