

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

Senior AI/LLM Engineer — GenAI-Powered Security Detection & Investigation Tooling

12 years building detection and investigation content for large-scale security operations, with hands-on experience putting GenAI/LLM tooling directly into live security workflows: prompt engineering for detection-content generation and false-positive analysis, LLM-driven orchestration of SIEM APIs, and reusable GenAI-powered tooling built for other detection engineers. Deep background integrating threat intelligence into detection and investigation content — directly relevant to SOC/threat-hunting investigation workflows — plus a strong Python and large-scale data engineering foundation (220+ ingested log sources, PySpark/GCP Dataproc, Elasticsearch).

CORE SKILLS

GenAI/LLM for Security: Prompt engineering for detection-content generation and false-positive/false-negative analysis; LLM-driven orchestration of SIEM APIs across 9 platforms for cross-customer detection deployment; reusable GenAI-powered tooling built for other engineers (e.g., automated cross-SIEM detection-rule conversion) — real production LLM-into-workflow experience

Threat Intelligence Integration: Integrating CTI (indicators, TTPs, actor/campaign context) into detection rule logic; enriching security alerts with threat intel for faster triage; using CTI directly during investigations and false-positive analysis (Forescout/Vedere Labs)

Python & Data Engineering: Python, SQL, pandas, NumPy, SciPy; data pipelines across 220+ ingested log sources; PySpark/SparkSQL on GCP Dataproc, Apache Beam/GCP Dataflow, BigQuery; Common Information Model / data standardization design

Detection Engineering & SIEM: Elasticsearch (queries, transforms, native detection rules, Beats, API), Splunk, MITRE ATT&CK-mapped detection content; GitLab CI/CD for detection-as-code with automated testing and staged/safe rollout

Cloud Platforms: AWS, GCP, Azure (Sentinel/Defender API orchestration); Docker; Kubernetes-orchestrated platform experience

PROFESSIONAL EXPERIENCE

Senior Cybersecurity Engineer — Shorepoint

Oct 2023 – Present

- Built production GenAI tooling for security automation: prompt engineering for false-positive triage and detection-content generation, LLM-orchestrated interaction with SIEM APIs across multiple platforms, and reusable GenAI-powered "skills" that convert detection rules across SIEM syntaxes for other engineers.
- Currently create and manage detection/alerting analytics (Splunk saved searches) directly supporting Treasury's Security Operations Center investigation workflows — translating threat-intel-informed and false-positive/false-negative investigation findings into refined detection logic (Treasury SOC / TSSOC, current project).
- Built an entirely new Elasticsearch-based security data platform from scratch for DOE/NNSA — ingesting CrowdStrike, Suricata, and Zeek telemetry, with a UEBA detection layer on custom data transforms, custom Kibana dashboards, and data-quality monitoring/alerting (DOE/NNSA Security Data Integration, completed).

Senior Threat Detection Engineer and Data Scientist — Forescout

Aug 2022 – Oct 2023

- Integrated threat intelligence from Forescout's in-house Vedere Labs research team directly into detection content and investigation workflows — using CTI to tune rule logic, enrich alerts for faster triage, and validate findings during false-positive analysis as a senior member of the data science and detection engineering team.

Threat Detection Engineer and Data Scientist — Trend Micro / Cysiv

Sep 2018 – Aug 2022

- Owned data engineering for a next-gen cloud SIEM at scale — pipelines for 220+ log sources, 50+ Logstash filters, and a Common Information Model standardizing fields across all of it — while creating and managing 2,300+ detection rules covering most of the MITRE ATT&CK matrix as a very early startup hire.
- Ran exploratory data analysis at scale (PySpark/SparkSQL on GCP Dataproc, Zeppelin notebooks) to validate new detection signal candidates before productionizing.

Security Data Scientist — Experian

Jan 2015 – Jan 2018

- Analyzed large-scale security log data to build custom detection models — DNS-based malware detection/mitigation and anomalous-behavior discovery across the network.

EDUCATION & CERTIFICATIONS

M.S. Physics — University of North Texas (2013) | B.S. Physics — Ball State University (2011)

Splunk User Certification · Splunk for Analytics and Data Science · Johns Hopkins (Coursera): R Programming, The Data Scientist's Toolbox

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

August 3, 2026

SOC/TH Team
Censys

Censys is building AI-powered investigation workflows that help analysts and threat hunters triage malicious infrastructure at internet scale — I've spent the last several years doing the adjacent version of exactly that: putting GenAI and LLM tooling directly into live security detection and investigation workflows, on top of a decade-plus of building the detection content those workflows run on.

At Shorepoint, I've built production GenAI tooling for security teams directly: prompt engineering for false-positive triage and automated detection-content generation, LLM-orchestrated interaction with SIEM APIs across nine platforms, and reusable GenAI-powered tooling that converts detection logic across SIEM syntaxes for other engineers. That's hands-on experience shaping how LLMs reason over security data and surface actionable findings — the same problem Censys is solving for internet-scale threat infrastructure.

I also bring direct SOC/investigation-workflow relevance: at Forescout, I integrated threat intelligence from the Vedere Labs research team into detection content and alert enrichment, using CTI to tune rule logic and validate findings during false-positive analysis — the same investigation-support work Censys' SOC-TH team builds tooling around. And at Trend Micro/Cysiv, I built the data engineering and detection infrastructure underneath all of it: pipelines for 220+ log sources, a Common Information Model standardizing the data, and 2,300+ detection rules across the MITRE ATT&CK matrix — the kind of internet-scale, high-signal data foundation that AI-powered investigation tooling needs to reason over.

I'd welcome the chance to talk through how that combination — production GenAI-for-security experience plus deep detection and investigation-content engineering — applies to shaping Censys' AI-powered SOC platform.

Kyle Main