

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

August 3, 2026

Data Science Platform Engineering Team
CrowdStrike

Before applying to help build data infrastructure for CrowdStrike's AI-driven security products, I built a security data platform that ingested CrowdStrike telemetry itself — a ground-up Elasticsearch-based ingestion and detection platform I built from scratch for a DOE/NNSA security program, pulling in CrowdStrike, Suricata, and Zeek data streams, then layering UEBA detection, data-quality monitoring, and custom dashboards on top of it.

That ground-up ownership runs through my whole career. I built the data engineering layer for 220+ heterogeneous log sources at a fast-growing security data company, including a Common Information Model that standardized schema and field types across every one of them — the same semantic-cataloging discipline large-scale AI/ML data platforms depend on. I've run large-scale distributed processing on GCP (Dataproc, PySpark/SparkSQL, BigQuery, Dataflow via a homegrown Apache Beam program), and I run every pipeline change through a full CI/CD process — a GitLab-based pipeline with automated testing, staged/safe rollout, and quality-metrics tracking before anything reaches production.

On the AI side, I've put GenAI directly into production security workflows: prompt engineering for detection-content generation and false-positive triage, and GenAI-driven orchestration of SIEM APIs across multiple platforms — real, hands-on experience shipping LLM-powered tooling into live operational use, which I'm eager to build on at the scale and depth this role calls for.

I'd welcome the chance to talk through how that combination — ground-up data pipeline ownership, CI/CD discipline, and hands-on GenAI tooling experience — applies to building CrowdStrike's next generation of AI-driven security data infrastructure.

Kyle Main