

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

Senior Data Engineer — Large-Scale Pipelines, GenAI/LLM Tooling & Platform Engineering

Data engineer with 12 years building large-scale data pipelines and detection-content platforms, including ground-up ingestion architecture for 220+ heterogeneous data sources, a Common Information Model standardizing schema across all of it, and Apache Beam/GCP Dataflow pipelines for high-volume historical retrieval. Direct, hands-on GenAI/LLM experience applying prompt engineering to production security workflows and orchestrating SIEM APIs via GenAI-powered tooling, plus a full CI/CD pipeline (GitLab) with automated testing, staged/safe rollout, and quality-metrics tracking — the same discipline this role's AI/ML data platforms require. Cybersecurity background includes hands-on ingestion and detection work built directly on CrowdStrike telemetry.

CORE SKILLS

Data Engineering & Pipeline Architecture: Ground-up ingestion pipelines for 220+ heterogeneous log/data sources; Common Information Model design — a data dictionary standardizing field names/types/schema across every source; 50+ Logstash filters for parsing/normalization; data-quality monitoring and alerting; homegrown Apache Beam program on GCP Dataflow for high-volume historical/cold-storage retrieval

Cloud & Distributed Data Processing: GCP (Dataproc, BigQuery, Dataflow); Spark/PySpark/SparkSQL for large-scale distributed processing; GCP serverless/event-driven enrichment; working familiarity with Kafka and Flink (hands-on exposure to a Flink-job environment)

GenAI / LLM Applications & Tooling: Prompt engineering for security use cases — data analysis, false-positive identification, automated detection-content generation; GenAI-driven orchestration of SIEM APIs across 9 platforms; reusable GenAI-powered tooling built for other engineers (e.g., automated cross-SIEM detection-rule conversion)

Multi-Platform Orchestration & CI/CD: Built and run a full CI/CD pipeline (GitLab) for detection/content deployment across 9 platforms via native APIs, with automated unit/integration testing, staged/safe production rollout, formally tracked quality metrics, and multithreaded parallel deployment

Platform & Access Engineering: Python, SQL, Git; Docker (custom images, containerized test environments); worked within a Kubernetes-orchestrated platform as a user; Cloud IAM policy/role implementation (AWS, GCP)

PROFESSIONAL EXPERIENCE

Senior Cybersecurity Engineer — Shorepoint

Oct 2023 – Present

- Built an entirely new data ingestion platform from scratch (DOE/NNSA project, completed) pulling CrowdStrike, Suricata, and Zeek telemetry into a central Elasticsearch environment — custom data transforms, a UEBA detection layer built on top of them, data-quality monitoring/alerting, and custom dashboards.
- Building production GenAI tooling for security automation (current): prompt engineering for detection-content generation and false-positive triage, plus GenAI-driven orchestration of SIEM APIs across multiple platforms.
- Supported data ingestion/quality engineering (CISA CDM at DOE, completed) across a combined Elasticsearch/Splunk environment; currently manages detection/alerting analytics (Splunk saved searches) for Treasury's SOC (TSSOC, current project).

Senior Threat Detection Engineer and Data Scientist — Forescout

Aug 2022 – Oct 2023

- Senior member of a data science and detection engineering team building signature, statistical, behavioral, and ML-based detection content and pipelines against massive-scale customer data on a cloud-based big-data platform.

Threat Detection Engineer and Data Scientist — Trend Micro / Cysiv

Sep 2018 – Aug 2022

- Owned data engineering for a next-gen cloud-based analytics platform at massive scale — 220+ ingested log sources, 50+ Logstash parsing/normalization filters, and a Common Information Model standardizing field names/types across all of it — as a very early hire building the pipeline from scratch.
- Built a homegrown Apache Beam program run via GCP Dataflow to fetch high-volume historical/cold-storage data on request; owned connector/pipeline health monitoring and troubleshooting.
- Ran exploratory data analysis at scale on GCP Dataproc using Zeppelin notebooks and PySpark/SparkSQL; created/managed 2,300+ detection rules covering most of the MITRE ATT&CK matrix, including ML-based clustering and time-series anomaly detection models.

Security Data Scientist — Experian

Jan 2015 – Jan 2018

- Analyzed large-scale security log data to build custom detection models — DNS-based malware detection/mitigation and anomalous-behavior discovery across the network.

EDUCATION & CERTIFICATIONS

M.S. Physics — University of North Texas (2013) | B.S. Physics — Ball State University (2011)

Splunk User Certification · Splunk for Analytics and Data Science · Johns Hopkins (Coursera): R Programming, The Data Scientist's Toolbox · Security Clearances: Top Secret (current), DOE Q (held), Public Trust (held)

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

August 3, 2026

Data Science Platform Engineering Team
CrowdStrike

Before applying to help build data infrastructure for CrowdStrike's AI-driven security products, I built a security data platform that ingested CrowdStrike telemetry itself — a ground-up Elasticsearch-based ingestion and detection platform I built from scratch for a DOE/NNSA security program, pulling in CrowdStrike, Suricata, and Zeek data streams, then layering UEBA detection, data-quality monitoring, and custom dashboards on top of it.

That ground-up ownership runs through my whole career. I built the data engineering layer for 220+ heterogeneous log sources at a fast-growing security data company, including a Common Information Model that standardized schema and field types across every one of them — the same semantic-cataloging discipline large-scale AI/ML data platforms depend on. I've run large-scale distributed processing on GCP (Dataproc, PySpark/SparkSQL, BigQuery, Dataflow via a homegrown Apache Beam program), and I run every pipeline change through a full CI/CD process — a GitLab-based pipeline with automated testing, staged/safe rollout, and quality-metrics tracking before anything reaches production.

On the AI side, I've put GenAI directly into production security workflows: prompt engineering for detection-content generation and false-positive triage, and GenAI-driven orchestration of SIEM APIs across multiple platforms — real, hands-on experience shipping LLM-powered tooling into live operational use, which I'm eager to build on at the scale and depth this role calls for.

I'd welcome the chance to talk through how that combination — ground-up data pipeline ownership, CI/CD discipline, and hands-on GenAI tooling experience — applies to building CrowdStrike's next generation of AI-driven security data infrastructure.

Kyle Main