

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

Senior Security Engineer — Data Infrastructure, Cloud IAM & Data Pipeline Engineering

Security engineer with 12 years of experience building the data pipelines, schema standards, and access-management frameworks that determine how data is ingested, modeled, and governed at scale. As an early hire at a cloud SIEM startup, built a Common Information Model — a data dictionary standardizing field names and types across 220+ ingested data sources — the same executable-schema discipline data classification and lineage propagation depend on. Hands-on Cloud IAM implementation in both AWS and GCP, plus API-level access governance (tokens, roles, permissions) across a dozen+ platforms. Comfortable writing production code and owning data infrastructure end to end, from architecture through operation.

CORE SKILLS

Data Modeling & Classification-Adjacent Engineering: Common Information Model (CIM) / data dictionary design — standardized field names and types across 220+ ingested log sources so downstream systems could act on schema automatically; data lineage awareness from designing and troubleshooting multi-source ingestion pipelines end to end

Cloud IAM & Access Governance: AWS and GCP IAM policy/role implementation; creation and governance of API tokens, roles, and permissions across 10+ security/data platforms as part of a production access-management framework

Data Pipeline Engineering: 220+ source data ingestion pipeline ownership; 50+ Logstash filters for parsing/normalization; GCP Dataflow/Apache Beam for large-scale historical data retrieval; PySpark on GCP Dataproc for exploratory analysis at scale; GCP serverless event-driven data enrichment

Cloud & Data Platforms: AWS, GCP, Azure (API orchestration); Elasticsearch (queries, transforms, Beats, API); Kafka/Flink exposure; Docker

Engineering & Delivery: Python, SQL, Git; GitLab CI/CD with automated testing and staged/safe rollout for pipeline and rule changes

PROFESSIONAL EXPERIENCE

Senior Cybersecurity Engineer — Shorepoint

Oct 2023 – Present

- Built an entirely new security data ingestion platform for DOE/NNSA from the ground up — architecting the pipeline that brought CrowdStrike, Suricata, and Zeek telemetry into a central Elasticsearch environment, then layering data-quality monitoring/alerting and a UEBA detection layer on top of custom data transforms (DOE/NNSA Security Data Integration project, completed).
- Supported data ingestion and data-quality efforts within an Elasticsearch/Splunk environment for DOE's Continuous Diagnostics and Mitigation (CDM) program, a continuous-monitoring initiative dependent on consistent, well-governed data models (CISA CDM at DOE, completed).
- Currently create and manage detection/alerting analytics (Splunk saved searches) supporting the Treasury Security Operations Center's incident response and case work, working directly with sensitive federal security data under strict access controls (Treasury SOC / TSSOC, current project).

Senior Threat Detection Engineer and Data Scientist — Forescout

Aug 2022 – Oct 2023

- Senior member of a data science and detection engineering team building signature, behavioral, statistical, and ML-based detection content against massive-scale customer telemetry on a cloud-based big-data platform.

Threat Detection Engineer and Data Scientist — Trend Micro / Cysiv

Sep 2018 – Aug 2022

- Designed and built a Common Information Model — a data dictionary standardizing field names and types across all parsed data — so every downstream system in a next-gen cloud SIEM could consume 220+ heterogeneous log sources through one consistent schema, the closest real analog in this history to executable, lineage-aware data classification.
- Owned data engineering for that ingestion pipeline at scale: built 50+ Logstash filters for parsing/normalization, deployed multiple Elasticsearch Beats variants for collection, and ran connector/collector health monitoring and troubleshooting to keep the pipeline reliable as a very early hire building the system from scratch.
- Built a homegrown Apache Beam program run on GCP Dataflow to fetch large volumes of historical cold-storage data for customers, and used GCP Dataproc/PySpark/SparkSQL for exploratory data analysis at scale.

Security Data Scientist — Experian

Jan 2015 – Jan 2018

- Analyzed large-scale security log data to build custom detection models — DNS-based malware detection/mitigation and anomalous-behavior discovery across the network.

EDUCATION & CERTIFICATIONS

M.S. Physics — University of North Texas (2013) | B.S. Physics — Ball State University (2011)

Splunk User Certification · Splunk for Analytics and Data Science · Johns Hopkins (Coursera): R Programming, The Data Scientist's Toolbox

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

August 3, 2026

Hiring Team
Grow Therapy

Grow Therapy's Senior Security Engineer role is fundamentally about making data classification, governance, and access control executable rather than aspirational — building the pipelines and standards that let every other system act on data sensitivity automatically. That's the exact problem I spent four years solving as an early hire building a cloud SIEM startup's data infrastructure from scratch.

At Trend Micro/Cysiv, I designed and built a Common Information Model — a data dictionary standardizing field names and types across 220+ heterogeneous log sources — so every downstream system in our next-gen SIEM could consume incoming data through one consistent, executable schema instead of guessing at its source by source. I owned that pipeline end to end: 50+ Logstash filters for parsing and normalization, multiple Elasticsearch Beats variants for collection, and a homegrown Apache Beam program on GCP Dataflow for large-scale historical data retrieval. That's the same discipline this role calls for — production code, real data pipelines, and a schema that the rest of the company can build on top of.

More recently at Shorepoint, I built DOE/NNSA's entire security data ingestion platform from a blank canvas — CrowdStrike, Suricata, and Zeek telemetry into a new Elasticsearch environment, with data-quality monitoring and alerting layered on top to keep the pipeline trustworthy over time. Alongside that, I've implemented Cloud IAM policies and roles in both AWS and GCP, and created and governed API tokens, roles, and permissions across a dozen-plus platforms as part of a production access-management framework — the same access-governance instinct this role's secure-data-access requirement is built around.

I'd welcome the chance to talk through how that data-infrastructure background applies to making the secure path the default one at Grow Therapy.

Kyle Main