

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

August 3, 2026

Hiring Team
OpenAI

I'd like to bring my data engineering background to the Applied team's work building the pipelines and canonical datasets that power OpenAI's analyses, safety systems, and product growth.

As one of the earliest engineers at Cysiv, a cybersecurity startup later acquired by Trend Micro and then Forescout, I built the data engineering foundation for our entire analytics platform from scratch: ingestion and pipelining for 220+ unique data sources, 50+ parsing/normalization filters, and a Common Information Model — a canonical data dictionary standardizing field names and types across every pipeline we ran. That last piece maps directly onto this role's charge to develop canonical datasets that track key metrics across a fast-growing product. I also built a homegrown Apache Beam program run on GCP Dataflow to retrieve large volumes of historical cold-storage data for customers, and ran exploratory data analysis at scale on GCP Dataproc clusters using PySpark and SparkSQL — the same distributed-processing muscle this role calls for with Spark.

Throughout my career, I've worked cross-functionally with infrastructure, data science, and product stakeholders to turn raw, high-volume event data into reliable, fault-tolerant systems that downstream teams could trust for decision-making — the same collaborative, data-integrity-first approach this role requires working across Infrastructure, Data Science, Product, Marketing, Finance, and Research. I'm drawn to the chance to apply that background to data that directly powers safety systems and the researchers building OpenAI's models.

I'd welcome the opportunity to talk through how this experience applies to building OpenAI's next data pipelines.

Kyle Main