Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

August 3, 2026

Hiring Team
OpenAI — Statsig / Core Experimentation

I've spent 12 years building the data pipelines and canonical data models that organizations depend on for trustworthy, large-scale analytics — exactly the foundation the Core Experimentation team needs underneath every experiment result it ships to the rest of OpenAI.

At Trend Micro/Cysiv, I owned data engineering for a next-gen cloud SIEM processing telemetry from 220+ unique log sources — building the ingestion pipelines, writing 50+ parsing/normalization filters, and designing a Common Information Model that standardized field names and types across every parsed data source, so downstream teams could query and analyze it consistently. I also built a homegrown Apache Beam program run via GCP Dataflow to retrieve large volumes of historical cold-storage data on demand, and used PySpark/SparkSQL on GCP Dataproc compute clusters for exploratory analysis at scale. That's the same shape of problem this role calls for: designing and managing pipelines that get raw event data cleanly into a data warehouse and turn it into canonical, trustworthy datasets.

Beyond the pipeline work, I bring a real statistics and data-science background to bear on that infrastructure — time-series anomaly detection, unsupervised clustering, and exploratory data analysis at scale, grounded in a quantitative M.S. Physics background in numerical data analysis and modeling. I understand what it takes to build data systems that hold up under statistical scrutiny, not just ones that move data from one place to another.

I'd welcome the chance to talk through how that background applies to building out Core Experimentation's data pipelines and canonical tables at OpenAI.

Kyle Main