

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

August 3, 2026

Hiring Team

OpenAI — Scaling Analytics, Industrial Compute

I've spent 12 years solving the exact problem underneath this role: connecting fragmented, heterogeneous systems into data that engineering and planning teams can actually trust and use — just applied to security telemetry instead of CPU and storage infrastructure. The underlying engineering is the same.

At Trend Micro/Cysiv, I owned data engineering for a next-gen cloud SIEM pulling in 220+ unique, heterogeneous log sources — designing a Common Information Model that gave every source a consistent representation of field names and types, and owning connector/collector health monitoring and troubleshooting to keep those integrations reliable in production. I also built reusable backend integrations against a dozen-plus SIEM and security platforms through their native REST APIs as part of a broader orchestration framework, including API token/role governance and multithreaded parallel execution — the same shape of work as connecting CPU/storage data across internal platforms, vendor APIs, and capacity systems. On the batch-processing side, I built a homegrown Apache Beam program run on GCP Dataflow to reliably retrieve large volumes of historical cold-storage data on demand, and used PySpark/SparkSQL on GCP Dataproc clusters for large-scale data analysis.

I'm comfortable navigating unfamiliar codebases and systems to figure out where a new integration should live, and I've built the data-quality and reconciliation monitoring that catches missing or inconsistent data before it reaches downstream consumers — direct experience with the reliability discipline this role needs applied to infrastructure data at scale.

I'd welcome the chance to talk through how that background applies to building out data integrations for Industrial Compute's CPU and storage platforms.

Kyle Main