

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

Data Engineer — Infrastructure Data Integrations, Pipelines & Fleet-Scale Systems

Data engineer with 12 years of experience connecting heterogeneous, fragmented systems into reliable, canonical data — the same core problem as unifying CPU, storage, and infrastructure telemetry across internal platforms and vendor systems. Built and ran ingestion integrations for 220+ unique log/data sources, owned connector/collector health monitoring and troubleshooting to keep integrations reliable in production, and designed a Common Information Model standardizing field names and types across every source. Built backend integrations against many third-party systems via native REST APIs, plus batch/scheduled data workloads with Apache Beam / GCP Dataflow, PySpark, and GCP Dataproc. Strong Python/SQL fundamentals with a track record of building data-quality and reconciliation checks that catch missing or inconsistent data.

CORE SKILLS

Data Engineering & Systems Integration: Ingestion pipeline design across 220+ heterogeneous data sources; connector/collector health monitoring and troubleshooting in production; Common Information Model / canonical data model design and standardization; data-quality and reconciliation monitoring/alerting to catch missing, stale, or inconsistent data; 50+ ETL/normalization filters (Logstash)

Backend & API Integration: Building integrations against third-party/vendor systems via native REST APIs across a dozen+ distinct platforms; reusable per-technology adapter design; API token/role/permission governance; multithreaded orchestration for parallel API-based workloads; GitLab CI/CD with automated testing and staged/safe production rollout

Batch, Distributed & Cloud Data Processing: Apache Beam / GCP Dataflow for large-scale batch data retrieval; PySpark / SparkSQL on GCP Dataproc compute clusters; BigQuery; cloud object storage (GCS buckets); GCP serverless/event-driven data enrichment; working exposure to Kafka and Flink

Engineering Fundamentals: Python, SQL, Git, GitLab CI/CD, automated testing, Docker; AWS, GCP, Azure

PROFESSIONAL EXPERIENCE

Senior Cybersecurity Engineer — Shorepoint

Oct 2023 – Present

- Built an entirely new data ingestion platform for DOE/NNSA from the ground up, connecting CrowdStrike, Suricata, and Zeek telemetry into a central Elasticsearch environment — data transforms, data-quality monitoring/alerting to catch inconsistent or missing data, and dashboards on top of it (DOE/NNSA Security Data Integration, completed).
- Supported data ingestion and data-quality engineering within an Elasticsearch/Splunk environment for DOE's Continuous Diagnostics and Mitigation (CDM) program, a large-scale federal continuous-monitoring data platform (completed).
- Currently builds and maintains analytics content (Splunk saved searches) supporting Treasury's Security Operations Center (Treasury SOC / TSSOC, current project).

Senior Threat Detection Engineer and Data Scientist — Forescout

Aug 2022 – Oct 2023

- Senior member of a data science and engineering team building detection/analytics content against massive-scale customer telemetry on cloud-based big-data platforms.

Threat Detection Engineer and Data Scientist — Trend Micro / Cysiv

Sep 2018 – Aug 2022

- Owned data engineering for a next-gen cloud-based SIEM at scale: built ingestion integrations for 220+ unique, heterogeneous log/data sources and designed a Common Information Model — a canonical data dictionary standardizing field names and types across every parsed source.
- Owned connector/collector health monitoring and troubleshooting to keep ingestion reliable in production; wrote 50+ Logstash parsing/normalization filters; assisted building "Loggify," a homegrown log parsing/filtering tool that replaced Logstash.
- Built a homegrown Apache Beam program run via GCP Dataflow to reliably fetch large volumes of historical cold-storage data on demand; ran exploratory data analysis at scale on GCP Dataproc compute clusters using PySpark/SparkSQL jobs loading data from cloud storage buckets.
- Built reusable per-technology backend integrations against a dozen+ SIEM/security platforms via native REST APIs as part of a broader orchestration framework, including API token/role/permission governance and multithreaded parallel orchestration.

Security Data Scientist — Experian

Jan 2015 – Jan 2018

- Analyzed large-scale security log datasets to build custom detection models — DNS-based malware detection/mitigation and anomalous-behavior discovery across network data.

EDUCATION & CERTIFICATIONS

M.S. Physics — University of North Texas (2013) | B.S. Physics — Ball State University (2011)

Splunk User Certification · Splunk for Analytics and Data Science · Johns Hopkins (Coursera): R Programming, The Data Scientist's Toolbox

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

August 3, 2026

Hiring Team

OpenAI — Scaling Analytics, Industrial Compute

I've spent 12 years solving the exact problem underneath this role: connecting fragmented, heterogeneous systems into data that engineering and planning teams can actually trust and use — just applied to security telemetry instead of CPU and storage infrastructure. The underlying engineering is the same.

At Trend Micro/Cysiv, I owned data engineering for a next-gen cloud SIEM pulling in 220+ unique, heterogeneous log sources — designing a Common Information Model that gave every source a consistent representation of field names and types, and owning connector/collector health monitoring and troubleshooting to keep those integrations reliable in production. I also built reusable backend integrations against a dozen-plus SIEM and security platforms through their native REST APIs as part of a broader orchestration framework, including API token/role governance and multithreaded parallel execution — the same shape of work as connecting CPU/storage data across internal platforms, vendor APIs, and capacity systems. On the batch-processing side, I built a homegrown Apache Beam program run on GCP Dataflow to reliably retrieve large volumes of historical cold-storage data on demand, and used PySpark/SparkSQL on GCP Dataproc clusters for large-scale data analysis.

I'm comfortable navigating unfamiliar codebases and systems to figure out where a new integration should live, and I've built the data-quality and reconciliation monitoring that catches missing or inconsistent data before it reaches downstream consumers — direct experience with the reliability discipline this role needs applied to infrastructure data at scale.

I'd welcome the chance to talk through how that background applies to building out data integrations for Industrial Compute's CPU and storage platforms.

Kyle Main