

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

Data Engineer — Data Pipelines, Data Quality & Operational Analytics at Scale

Data engineer with 12 years of experience building the pipelines, canonical data models, and data-quality systems that turn fragmented operational data into trusted analytics — the same foundation Scaling Analytics needs to support infrastructure deployment, capacity planning, and operational decision-making. Built data ingestion and pipelining for 220+ unique data sources as an early startup hire, including a Common Information Model standardizing field names and types across every source. Deep hands-on experience with PySpark/GCP Dataproc for distributed processing and BigQuery for analytical data warehousing, plus a homegrown Apache Beam program on GCP Dataflow for large-scale batch data retrieval. Strong SQL/Python fundamentals with a track record of building data-quality monitoring and dashboards that stakeholders across engineering, operations, and leadership rely on for decisions.

CORE SKILLS

Data Engineering & Pipelines: Production data pipeline design and operation across 220+ heterogeneous sources; Common Information Model / canonical data model design and standardization; ETL/normalization (50+ Logstash filters); data-quality checks, monitoring, and alerting in production; connector/collector health monitoring and troubleshooting

Distributed & Cloud Data Processing: PySpark / SparkSQL on GCP Dataproc compute clusters, BigQuery, Apache Beam on GCP Dataflow for batch/historical data retrieval, GCS object storage, GCP serverless/event-driven data enrichment

Programming & Analysis: Python, SQL, Git, GitLab CI/CD with automated testing and staged/safe rollout; Pandas, NumPy, scikit-learn

Reporting, Dashboards & Cross-Functional Delivery: Kibana dashboard design for operational decision-making; translating cross-functional stakeholder requirements (security, data science, SOC operations) into technical data solutions; sustained delivery in fast-paced, evolving-priority environments

Cloud Platforms: GCP, AWS, Azure

PROFESSIONAL EXPERIENCE

Senior Cybersecurity Engineer — Shorepoint

Oct 2023 – Present

- Built an entirely new data ingestion platform from the ground up, pipelining three high-volume operational telemetry sources (CrowdStrike, Suricata, Zeek) into a central Elasticsearch environment — data transforms, data-quality monitoring/alerting, and custom dashboards that gave stakeholders real-time visibility into system health (DOE/NSA Security Data Integration, completed).
- Currently builds and maintains analytics content that turns raw operational event data into decision-ready outputs supporting a security operations center's case and incident workflows (Treasury SOC / TSSOC, current project).
- Supported data ingestion and data-quality engineering within an Elasticsearch/Splunk environment for a large-scale continuous-monitoring data platform (CISA CDM at DOE, completed).

Senior Threat Detection Engineer and Data Scientist — Forescout

Aug 2022 – Oct 2023

- Senior member of a data science and engineering team building analytics content against massive-scale customer data on cloud-based big-data platforms — signature, statistical, time-series, and ML-based analysis feeding operational decisions.

Threat Detection Engineer and Data Scientist — Trend Micro / Cysiv

Sep 2018 – Aug 2022

- Owned data engineering for a next-gen cloud-based analytics platform at scale: built ingestion pipelines for 220+ unique, heterogeneous data sources and designed a Common Information Model — a canonical data dictionary standardizing field names and types across every source, enabling consistent reporting and analysis.
- Wrote 50+ Logstash parsing/normalization filters and owned connector/collector health monitoring and troubleshooting to keep pipelines reliable in production; assisted building "Loggify," a homegrown log parsing/filtering tool that replaced Logstash.
- Built a homegrown Apache Beam program run via GCP Dataflow to reliably retrieve large volumes of historical/cold-storage data on demand; ran exploratory data analysis at scale on GCP Dataproc compute clusters using PySpark/SparkSQL jobs against cloud storage buckets.
- Built data-quality monitoring and alerting content to catch missing or inconsistent data across the pipeline, and created custom dashboards translating raw data into metrics used by analysts and leadership for day-to-day decisions.

Security Data Scientist — Experian

Jan 2015 – Jan 2018

- Analyzed large-scale operational log datasets to build custom detection models — DNS-based malware detection/mitigation and anomalous-behavior discovery across network data.

EDUCATION & CERTIFICATIONS

M.S. Physics — University of North Texas (2013) | B.S. Physics — Ball State University (2011)

Splunk User Certification · Splunk for Analytics and Data Science · Johns Hopkins (Coursera): R Programming, The Data Scientist's Toolbox

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

August 3, 2026

Hiring Team

OpenAI — Scaling Analytics

I've spent 12 years building the pipelines and canonical data models that turn fragmented, high-volume operational data into analytics people actually trust — the same discipline Scaling Analytics needs to give OpenAI's infrastructure organization visibility into hardware, capacity, and site operations at global scale.

At Trend Micro/Cysiv, I owned data engineering for a next-gen analytics platform pulling in 220+ unique, heterogeneous data sources — designing a Common Information Model that gave every source a consistent schema for reporting and analysis, writing 50+ normalization filters, and owning connector/collector health monitoring so the pipelines stayed reliable in production. On the distributed-processing side, I built a homegrown Apache Beam program run on GCP Dataflow to reliably retrieve large volumes of historical data on demand, and used PySpark/SparkSQL on GCP Dataproc clusters alongside BigQuery for large-scale analysis. More recently at Shorepoint, I built a data ingestion and analytics platform from scratch for a federal client — pipelines, data-quality monitoring, and dashboards that gave operators real-time visibility into system health, the same shape of work as building trusted datasets for hardware and capacity operations.

Across every role, I've translated requirements from stakeholders outside my own team — security analysts, SOC operators, data scientists — into pipelines and dashboards they could rely on for daily decisions, and I've built the data-quality checks that catch problems before they reach those decisions. I'm comfortable operating in fast-moving, ambiguous environments — I was an early hire at Cysiv, building core data infrastructure as the company scaled from a small team into an established SIEM provider.

I'd welcome the chance to talk through how that background applies to building out the analytical foundations for OpenAI's infrastructure organization.

Kyle Main