

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

Principal Data Engineer — Elasticsearch-Native Platforms & Large-Scale Data Pipelines

Data engineer with 12 years building and operating large-scale data pipelines and analytics platforms, anchored by deep, cross-employer hands-on Elasticsearch experience — from ES queries and transforms to native detection rules and API-level cluster work. Built the data engineering layer for 220+ heterogeneous data sources from the ground up, including a Common Information Model standardizing schema across every source, plus GCP Dataflow/Apache Beam pipelines for large-scale historical data retrieval. Strong Python foundation with statistical/ML technique experience, real team-lead/sprint-lead history mentoring peers on pipeline and detection design, and direct experience integrating threat intelligence into data/detection content.

CORE SKILLS

Elasticsearch & Search-Based Platforms: ES queries (Query DSL), ES transforms, Logstash (50+ custom filters), ES Beats (multiple log-collection variants), native ES detection rules/alerting, ES API (cluster/index management), Kibana dashboarding — core platform across three consecutive employers

Data Engineering at Scale: 220+ unique source ingestion pipelines, Apache Beam / GCP Dataflow for historical & cold-storage retrieval, Common Information Model / data-dictionary design, data-quality monitoring & alerting, staged/safe production rollout

Python & Analytics: Python, SQL, statistical & ML techniques (clustering, time-series anomaly detection), PySpark / SparkSQL / GCP Dataproc / BigQuery for large-scale distributed processing

Cloud, Streaming & Orchestration: AWS / GCP / Azure; working familiarity with Kafka and Flink; hands-on experience within a Kubernetes-orchestrated platform; GitLab CI/CD pipeline orchestration with automated testing and staged rollout

AI/LLM Applications & Threat Intelligence: Prompt engineering and GenAI-powered tooling for detection/data-analysis workflows; extensive experience integrating threat intelligence (CTI) into detection and data-analytics content

PROFESSIONAL EXPERIENCE

Senior Cybersecurity Engineer — Shorepoint

Oct 2023 – Present

- Built an entirely new Elasticsearch-based data platform (DOE/NNSA project, completed) ingesting CrowdStrike, Suricata, and Zeek data streams — designed ES transforms powering a UEBA detection layer, custom Kibana dashboards, and data-quality alerting.
- Supported data ingestion and data-quality engineering (CISA CDM project, completed) across a combined Elasticsearch/Splunk environment for a large federal data platform.
- Currently builds and maintains detection/analytics content against a large-scale Splunk data environment for Treasury's SOC (current project).

Senior Threat Detection Engineer and Data Scientist — Forescout

Aug 2022 – Oct 2023

- Senior member of a data science/detection engineering team building pipelines and analytics content — signature, statistical, behavioral, and ML-based — against massive-scale customer data on cloud-based big-data tooling.
- Integrated threat intelligence (Vedere Labs CTI) into detection logic and alert enrichment, using actor/campaign context to speed triage and tune rule precision.

Threat Detection Engineer and Data Scientist — Trend Micro / Cysiv

Sep 2018 – Aug 2022

- Early hire on a next-gen SIEM built natively on Elasticsearch/Kibana — wrote and managed detection rules and queries directly against ES indexes as core detection content, not just log shipping.
- Owned data engineering/pipelining for 220+ unique log sources; built 50+ Logstash filters for parsing/normalization and authored a Common Information Model — a data dictionary standardizing field names/types across every parsed source.
- Built a homegrown Apache Beam program run via GCP Dataflow to fetch high-volume historical cold-storage data on request; owned connector/collector health monitoring.
- Served in a team-lead capacity running sprints and mentoring detection engineers on ES query/transform design, rule development, and pipeline architecture.

Security Data Scientist — Experian

Jan 2015 – Jan 2018

- Analyzed large-scale security datasets to build custom models for anomalous-behavior detection; built DNS-based malware detection/mitigation logic on production data.

EDUCATION & CERTIFICATIONS

M.S. Physics — University of North Texas (2013) | B.S. Physics — Ball State University (2011)

Splunk User Certification · Splunk for Analytics and Data Science · Johns Hopkins (Coursera): R Programming, The Data Scientist's Toolbox

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

August 3, 2026

Digital Risk Protection Hiring Team
Recorded Future

Data convergence across disparate risk data streams is a problem I've solved from the ground up before: as an early hire building a next-gen SIEM, I owned the data engineering layer for 220+ distinct, messy sources, including a Common Information Model that standardized field names and types across every one of them so every downstream consumer — dashboards, detection rules, analytics — could build on one consistent schema instead of managing 220 different ones. That's the exact discipline aligning and ingesting Malicious Sites, Identity, and Surface Web data into a unified Digital Risk Protection system and Intelligence Graph requires.

Elasticsearch has been the core platform underneath nearly all of that work — across three consecutive employers I've written and tuned ES queries and transforms, deployed Beats for collection, built native ES detection rules and Kibana dashboards, and worked directly against the ES API for cluster and index management, not just through a UI. I've paired that with production GCP pipeline work (Dataproc, Dataflow via a homegrown Apache Beam program, PySpark/SparkSQL) and a strong Python foundation for the statistical and ML techniques this role calls for. I've also run point on integrating threat intelligence directly into detection and analytics content — using actor and campaign context to tune rule logic and speed triage — which maps directly onto the kind of digital risk data Recorded Future's Intelligence Graph is built to explain.

I'd welcome the chance to bring that pipeline-ownership, Elasticsearch depth, and threat-intelligence background to Recorded Future's Digital Risk Protection team.

Kyle Main