

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

August 3, 2026

Hiring Team

Upstart

Upstart's Data Security program is built on a principle I've spent years working from the access-control side: least-privilege only works if the systems governing it are engineered, not improvised. That's the problem I've been solving across identity/access governance and data pipeline engineering for the past several years.

I've implemented Cloud IAM policies and roles in both AWS and GCP, and — as part of building a multi-SIEM detection-as-code orchestration framework — created and governed API tokens, roles, and permissions across a dozen-plus security platforms (Microsoft Sentinel, Microsoft Defender, Google SecOps, Splunk, CrowdStrike, SentinelOne, Sumo Logic, Palo Alto XSIAM, Devo, and more), all via their native APIs. That framework runs on a GitLab CI/CD pipeline I built with automated unit/integration tests, tracked rule-quality metrics, and staged/safe rollout before anything reaches production — the same engineering discipline that turns a least-privilege access model from policy into something real and repeatable.

On the data side, I was an early hire building a cloud SIEM startup's data infrastructure from scratch: I designed a Common Information Model standardizing field names and types across 220+ ingested log sources, owned 50+ Logstash parsing/normalization filters, and built a homegrown Apache Beam program on GCP Dataflow for large-scale historical data retrieval. More recently at Shorepoint, I built DOE/NNSA's entire security data ingestion platform from a blank canvas, with data-quality monitoring and alerting layered on top to keep it trustworthy over time. That combination — access governance engineering plus hands-on data pipeline ownership — is exactly the intersection this role sits at.

I'd welcome the chance to talk through how that background applies to building Upstart's data security program.

Kyle Main