

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

Senior Security Engineer — Data Security, Cloud IAM & Access Governance

Security engineer with 12 years of experience building the identity, access-governance, and data pipeline systems that determine who and what can reach sensitive data. Hands-on Cloud IAM policy/role implementation in both AWS and GCP, plus API-level access governance — creating and managing tokens, roles, and permissions across a dozen-plus security platforms as part of a production orchestration framework built around least-privilege principles. Deep data-engineering background (220+ ingested source pipelines, schema standardization, large-scale cloud data processing) gives real technical grounding for data protection work across diverse data domains, backed by production Python engineering and full CI/CD ownership.

CORE SKILLS

Cloud IAM & Least-Privilege Access Governance: AWS and GCP IAM policy/role implementation; creation and governance of API tokens, roles, and permissions across 10+ security platforms as part of a production access-management framework built around least-privilege principles

API-Level Access Management & Multi-Platform Orchestration: Built per-platform API adapters and orchestration logic across Microsoft Sentinel, Microsoft Defender, Google SecOps (Chronicle), Splunk, CrowdStrike, SentinelOne, Sumo Logic, Palo Alto XSIAM, Devo, and ArcSight — real hands-on identity/access and permission-scoping work at the API level, not just rule management

Data Engineering & Pipeline Building: Designed and built a Common Information Model (data dictionary) standardizing field names/types across 220+ ingested log sources; 50+ Logstash filters; multiple Elasticsearch Beats variants; GCP Dataflow/Apache Beam for large-scale historical data retrieval; PySpark on GCP Dataproc for analysis at scale

Data Security & Detection Platforms: Elasticsearch (queries, transforms, Beats, native detection rules, API); Splunk; data-quality monitoring and alerting for federal-scale ingestion platforms; security detection content across signature, statistical, behavioral, and ML-based rule types

Engineering & Delivery: Python, SQL, Git, GitLab CI/CD pipeline for detection-as-code with automated unit/integration tests, tracked rule-quality metrics, and staged/safe production rollout; comfortable Docker user

PROFESSIONAL EXPERIENCE

Senior Cybersecurity Engineer — Shorepoint

Oct 2023 – Present

- Currently create and manage detection/alerting analytics (Splunk saved searches) directly supporting the Treasury Security Operations Center's incident response and case work, operating under strict least-privilege access controls on sensitive federal data (Treasury SOC / TSSOC, current project).
- Built an entirely new security data ingestion platform for DOE/NNSA from the ground up — CrowdStrike, Suricata, and Zeek telemetry into a central Elasticsearch environment — including a UEBA detection layer on custom data transforms, data-quality monitoring/alerting content, and custom dashboards (DOE/NNSA Security Data Integration project, completed).
- Supported data ingestion and data-quality efforts within an Elasticsearch/Splunk environment for DOE's Continuous Diagnostics and Mitigation (CDM) program (completed).

Senior Threat Detection Engineer and Data Scientist — Forescout

Aug 2022 – Oct 2023

- Senior member of a data science and detection engineering team building signature, behavioral, statistical, and ML-based detection content against massive-scale customer telemetry on a cloud-based big-data platform.

Threat Detection Engineer and Data Scientist — Trend Micro / Cysiv

Sep 2018 – Aug 2022

- Designed and built a Common Information Model — a data dictionary standardizing field names and types across all parsed data — so every downstream system in a next-gen cloud SIEM could consume 220+ heterogeneous log sources through one consistent schema, as a very early hire building the data infrastructure and rules engine from scratch.
- Owned data engineering for that ingestion pipeline at scale: built 50+ Logstash filters for parsing/normalization, deployed multiple Elasticsearch Beats variants for collection, and ran connector/collector health monitoring and troubleshooting to keep the pipeline reliable.
- Created and managed 2,300+ individual detection rules covering most of the MITRE ATT&CK matrix, and built a homegrown Apache Beam program on GCP Dataflow for large-scale historical cold-storage data retrieval.

Security Data Scientist — Experian

Jan 2015 – Jan 2018

- Analyzed large-scale security log data to build custom detection models — DNS-based malware detection/mitigation and anomalous-behavior discovery across the network.

EDUCATION & CERTIFICATIONS

M.S. Physics — University of North Texas (2013) | B.S. Physics — Ball State University (2011)

Splunk User Certification · Splunk for Analytics and Data Science · Johns Hopkins (Coursera): R Programming, The Data Scientist's Toolbox

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

August 3, 2026

Hiring Team

Upstart

Upstart's Data Security program is built on a principle I've spent years working from the access-control side: least-privilege only works if the systems governing it are engineered, not improvised. That's the problem I've been solving across identity/access governance and data pipeline engineering for the past several years.

I've implemented Cloud IAM policies and roles in both AWS and GCP, and — as part of building a multi-SIEM detection-as-code orchestration framework — created and governed API tokens, roles, and permissions across a dozen-plus security platforms (Microsoft Sentinel, Microsoft Defender, Google SecOps, Splunk, CrowdStrike, SentinelOne, Sumo Logic, Palo Alto XSIAM, Devo, and more), all via their native APIs. That framework runs on a GitLab CI/CD pipeline I built with automated unit/integration tests, tracked rule-quality metrics, and staged/safe rollout before anything reaches production — the same engineering discipline that turns a least-privilege access model from policy into something real and repeatable.

On the data side, I was an early hire building a cloud SIEM startup's data infrastructure from scratch: I designed a Common Information Model standardizing field names and types across 220+ ingested log sources, owned 50+ Logstash parsing/normalization filters, and built a homegrown Apache Beam program on GCP Dataflow for large-scale historical data retrieval. More recently at Shorepoint, I built DOE/NNSA's entire security data ingestion platform from a blank canvas, with data-quality monitoring and alerting layered on top to keep it trustworthy over time. That combination — access governance engineering plus hands-on data pipeline ownership — is exactly the intersection this role sits at.

I'd welcome the chance to talk through how that background applies to building Upstart's data security program.

Kyle Main