

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

Senior Detection & Response Engineer — GenAI for Defensive Security, Detection-as-Code

Senior detection and response engineer with 12 years building the detection content and defensive tooling that protect real production environments, plus genuine hands-on experience applying generative AI directly to defensive security workflows. Built the detection rules engine from scratch at a next-gen cloud SIEM startup (Cysiv, spun out of Trend Micro, later acquired by Forescout), creating 2,300+ MITRE ATT&CK-covering rules. Currently support live incident/case response for a federal SOC. Applied GenAI directly to detection engineering — prompt-driven false-positive triage, new detection-content generation, and automated rule translation across SIEM syntaxes — the kind of AI-for-defensive-security work this role's mandate to safely advance model capability in secure coding and vulnerability remediation calls for.

CORE SKILLS

Detection & Response: 2,300+ detection rules covering most of the MITRE ATT&CK matrix; current incident/case response for a federal SOC; formal detection-quality metrics (coverage, precision/false-positive rate) with staged/safe rollout before production

GenAI / LLM Applied to Defensive Security: Prompt engineering to analyze security data, identify false positives, and generate new detection content; GenAI-driven SIEM API orchestration across customers/platforms; built reusable GenAI-powered tooling that automates detection-rule translation between SIEM rule syntaxes — hands-on experience directing model output toward real defensive-security outcomes and validating it against analyst judgment before trusting it

Multi-SIEM Detection-as-Code & Orchestration: Rule/content orchestration via native APIs across nine SIEM platforms (Microsoft Sentinel, Microsoft Defender, Google SecOps, Splunk, CrowdStrike, SentinelOne, Sumo Logic, Palo Alto XSIAM, Devo); GitLab CI/CD with automated tests and multithreaded parallel deployment

Vulnerability & Data Engineering: Ingested Tenable vulnerability scan data into a SIEM platform and built analytics/detection content on top of it; 220+ log source data engineering; Common Information Model design

Cloud Security Platform Engineering: Built a cyber-defense platform from the ground up on Elasticsearch — CrowdStrike (EDR), Suricata, and Zeek telemetry, UEBA detection layer, data-quality monitoring — for a DOE/NNSA cloud-adjacent security program

Engineering: Python, SQL, PyTorch/scikit-learn, Git, GitLab CI/CD, Docker

PROFESSIONAL EXPERIENCE

Senior Cybersecurity Engineer — Shorepoint

Oct 2023 – Present

- Currently create and manage detection/alerting analytics directly supporting a federal SOC's live incident response and case work (Treasury SOC / TSSOC, current project).
- Built an entirely new cyber-defense platform from the ground up — CrowdStrike (EDR), Suricata, and Zeek telemetry into a central Elasticsearch environment — including a UEBA detection layer and data-quality monitoring/alerting content (DOE/NNSA, completed).

Senior Threat Detection Engineer and Data Scientist — Forescout

Aug 2022 – Oct 2023

- Senior member of a data science and detection engineering team building signature, behavioral, statistical, and ML-based detection content against massive-scale customer telemetry, incorporating threat intel from Forescout's in-house Vedere Labs research team.

Threat Detection Engineer and Data Scientist — Trend Micro / Cysiv

Sep 2018 – Aug 2022

- Very early hire at a next-gen cloud SIEM startup — built the detection rules engine from scratch and created/managed 2,300+ individual detection rules covering most of the MITRE ATT&CK matrix, plus 50+ data filters, against 220+ ingested log sources.
- Applied GenAI directly to detection engineering: prompt engineering for false-positive identification and new rule generation, and reusable GenAI-powered tooling that automates converting detection rules between SIEM rule syntaxes.
- Ran exploratory data analysis at scale (GCP Dataproc, PySpark/SparkSQL) using clustering and time-series anomaly detection to develop new detection content.

Security Data Scientist — Experian

Jan 2015 – Jan 2018

- Analyzed large-scale security log data to build custom detection models — DNS-based malware detection/mitigation and anomalous-behavior discovery across the network.

EDUCATION & CERTIFICATIONS

M.S. Physics — University of North Texas (2013) | B.S. Physics — Ball State University (2011)

Splunk User Certification · Splunk for Analytics and Data Science · Johns Hopkins (Coursera): R Programming, The Data Scientist's Toolbox

Kyle Main

Dallas/Ft. Worth Area | 469-545-3791 | main.kyle87@gmail.com

August 7, 2026

Hiring Team

Anthropic — Horizons, Cybersecurity RL

I'm a detection and response engineer who has spent 12 years building defensive security content, and for the last several years I've been applying generative AI directly to that work — which is exactly the intersection this role sits at. I'd like to bring that background to the Cybersecurity RL team.

As a very early hire at Cysiv (a next-gen cloud SIEM startup that spun out of Trend Micro, later acquired by Forescout), I built the detection rules engine from scratch and created and managed 2,300+ individual detection rules covering most of the MITRE ATT&CK matrix. On top of that core detection-engineering work, I've built real, hands-on GenAI tooling for defensive security: prompt engineering to analyze security data and identify false positives, GenAI-driven orchestration of detection content across nine SIEM platforms, and reusable GenAI-powered tooling that automates translating detection rules between SIEM syntaxes. I currently support live incident and case response for a federal SOC, so I bring the applied defensive-work perspective this role is built around — someone who has lived inside detection and response and is now curious about how models can meaningfully augment that work.

I've also built detection content directly on top of ingested vulnerability-scan data (Tenable), giving me real working familiarity with the secure-coding and vulnerability-remediation problem space this team is advancing model capability in — not just detection after the fact, but understanding where the underlying weaknesses come from.

I'd welcome the chance to talk through how that domain expertise could support the Cybersecurity RL team's work advancing Claude's capabilities in secure coding and vulnerability remediation.

Kyle Main